

# Machine Learning 1


Dora Deng

2025-01-28

## First up kmeans()

Demo of using kmeans() function in base R. First make up some data with a known structure.

```
tmp <- c(rnorm(30, -3), rnorm(30, 3))
x <- cbind(x = tmp, y = rev(tmp))
plot(x)
```




Now we have some made up data in x let's see how kmeans works with this data.

```
k <- kmeans(x, centers = 2, nstart = 20)
k
```

```
## K-means clustering with 2 clusters of sizes 30, 30
##
## Cluster means:
```





## Now for Hierarchical Clustering - `hclust()`

We will cluster the same data `x` with the `hclust()`. In this case `hclust()` requires a distance matrix as input.


```
hc <- hclust(dist(x))
hc

##
## Call:
## hclust(d = dist(x))
##
## Cluster method   : complete
## Distance        : euclidean
## Number of objects: 60
```

Let's plot our `hclust` result

```
plot(hc)
```

## Cluster Dendrogram



```
dist(x)
hclust (*, "complete")
```


To get our cluster membership vector we need to “cut” the tree with the `cutree()`

```
grps <- cutree(hc, h = 8)
grps
```

```
## [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
## [39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

Now plot our data with the `hclust()` results.

```
plot(x, col = grps)
```



## Principal Component Analysis (PCA)

### PCA of UK food data


Read data from webstie and try a few visualizations.

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url, row.names = 1)
x
```


##	England	Wales	Scotland	N.Ireland
## Cheese	105	103	103	66
## Carcass_meat	245	227	242	267
## Other_meat	685	803	750	586
## Fish	147	160	122	93
## Fats_and_oils	193	235	184	209
## Sugars	156	175	147	139
## Fresh_potatoes	720	874	566	1033
## Fresh_Veg	253	265	171	143
## Other_Veg	488	570	418	355
## Processed_potatoes	198	203	220	187
## Processed_Veg	360	365	337	334
## Fresh_fruit	1102	1137	957	674
## Cereals	1472	1582	1462	1494
## Beverages	57	73	53	47
## Soft_drinks	1374	1256	1572	1506

```
## Alcoholic_drinks      375   475   458   135
## Confectionery         54    64    62    41

cols <- rainbow(nrow(x))
barplot(as.matrix(x), col = cols, beside = TRUE)
```



```
pairs(x, col = cols)
```



PCA to the rescue!!

The main base R PCA function is called `prcomp()` and we will need to give it the transpose of our input data!

```
#t(x)
pca <- prcomp( t(x))
pca
```

```
## Standard deviations (1, ..., p=4):
## [1] 3.241502e+02 2.127478e+02 7.387622e+01 2.921348e-14
##
## Rotation (n x k) = (17 x 4):
##
##          PC1          PC2          PC3          PC4
## Cheese      -0.056955380  0.016012850  0.02394295 -0.409382587
## Carcass_meat  0.047927628  0.013915823  0.06367111  0.729481922
## Other_meat   -0.258916658 -0.015331138 -0.55384854  0.331001134
## Fish        -0.084414983 -0.050754947  0.03906481  0.022375878
## Fats_and_oils -0.005193623 -0.095388656 -0.12522257  0.034512161
## Sugars       -0.037620983 -0.043021699 -0.03605745  0.024943337
## Fresh_potatoes  0.401402060 -0.715017078 -0.20668248  0.021396007
## Fresh_Veg    -0.151849942 -0.144900268  0.21382237  0.001606882
## Other_Veg    -0.243593729 -0.225450923 -0.05332841  0.031153231
## Processed_potatoes -0.026886233  0.042850761 -0.07364902 -0.017379680
## Processed_Veg -0.036488269 -0.045451802  0.05289191  0.021250980
## Fresh_fruit  -0.632640898 -0.177740743  0.40012865  0.227657348
## Cereals      -0.047702858 -0.212599678 -0.35884921  0.100043319
## Beverages    -0.026187756 -0.030560542 -0.04135860 -0.018382072
```


```
## Soft_drinks      0.232244140  0.555124311 -0.16942648  0.222319484
## Alcoholic_drinks -0.463968168  0.113536523 -0.49858320 -0.273126013
## Confectionery    -0.029650201  0.005949921 -0.05232164  0.001890737
```

```
attributes(pca)
```

```
## $names
## [1] "sdev"      "rotation" "center"    "scale"     "x"
##
## $class
## [1] "prcomp"
```

To make our new PCA plot (a.k.a. PCA score plot) we access `pca$x`


```
plot(pca$x[,1], pca$x[,2])
text(pca$x[,1], pca$x[,2], colnames(x))
```



color up the plot

```
country_cols <- c("orange", "red", "blue", "green")
plot(pca$x[,1], pca$x[,2])
text(pca$x[,1], pca$x[,2], colnames(x), col = country_cols)
```





## PCA of RNA-seq data:

Read in data from website:

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

```
##          wt1 wt2  wt3  wt4 wt5 ko1 ko2 ko3 ko4 ko5
## gene1  439 458  408  429 420  90  88  86  90  93
## gene2  219 200  204  210 187 427 423 434 433 426
## gene3 1006 989 1030 1017 973 252 237 238 226 210
## gene4  783 792  829  856 760 849 856 835 885 894
## gene5  181 249  204  244 225 277 305 272 270 279
## gene6  460 502  491  491 493 612 594 577 618 638
```

Q.10: How many genes and samples are in this data set?

```
pca <- prcomp(t(rna.data))
```

```
# There is a nice summary of how well PCA is doing
summary(pca)
```


```
## Importance of components:
```

```
##          PC1      PC2      PC3      PC4      PC5      PC6
## Standard deviation 2214.2633 88.9209 84.33908 77.74094 69.66341 67.78516
## Proportion of Variance 0.9917 0.0016 0.00144 0.00122 0.00098 0.00093
## Cumulative Proportion 0.9917 0.9933 0.99471 0.99593 0.99691 0.99784
```

```
##
##          PC7      PC8      PC9      PC10
## Standard deviation 65.29428 59.90981 53.20803 2.662e-13
## Proportion of Variance 0.00086 0.00073 0.00057 0.000e+00
## Cumulative Proportion 0.99870 0.99943 1.00000 1.000e+00
```

Do our PCA plot of this RNA-Seq data

```
plot(pca$x[,1], pca$x[,2], xlab = "PC1", ylab = "PC2")
```



```
plot(pca$x[,1], pca$x[,2], xlab = "PC1", ylab = "PC2")
text(pca$x[,1], pca$x[,2], colnames(rna.data))
```

